





LLM Agents


Егор Спирин



1. Заставить модель рассуждать
2. Исправлять ошибки за собой
3. Использовать внешние инструменты
4. Взаимодействие с другими БЯМ



Quick Recap




Языковые модели

Умеем предсказывать следующий токен 🖐️

Модель обучается в несколько стадий:

1. Pre-Train — много текста, "впитывание" знаний
2. Instruction Tuning — добавляем инструкции, "пользуемся" знаниями
3. Preference Optimization — выравниваем ответы, "нравимся" пользователю



GPT-2, 3, ...

"Our largest model, GPT-2, is a 1.5B parameter Transformer that achieves state of the art results on 7 out of 8 tested language modeling datasets in a zero-shot setting but still underfits WebText"

Модель обученная на предсказывание следующего слова оказалась SOTA:

- 👉 POS-tagging
- 👉 Summarization
- 👉 Translation
- 👉 ...

И все это без обучения под задачи против специально обученных моделей 🔥

Prompt Crafting



Prompt a.k.a. Затравка

Текстовый запрос или инструкция, подаваемая языковой модели для генерации ответа или выполнения задачи

👉 Управление моделью

👉 Передача знаний

👉 Позволяют раскрыть весь потенциал БЯМ

Размер контекста ограничен моделью $\Rightarrow$ размер промпта также ограничен

Question: how many apples in the bucket?

Answer: there are 52 apples.

Few-shot

Можем показать модели несколько примеров с ожидаемым поведением

Примеры с
паттерном

Question: Elon Musk
Answer: nk

Question: Bill Gates
Answer: ls

Сама задача

Question: Sam Altman
Answer: ??

Человек легко ответить — он умеет размышлять (reasoning), в отличие от LLM 🤖

Chain-of-Thought

Давайте добавим шаг с рассуждением!

Question: Elon Musk

Answer: the last letter of "Elon" is "n". The last letter of "Musk" is "k". Concatenating "n", "k" leads to "nk". So the output is "nk".

Question: Bill Gates

Answer: the last letter of "Bill" is "l". The last letter of "Gates" is "s". Concatenating "l", "s" leads to "ls". So the output is "ls".

Question: Sam Altman

Answer: the last letter of "Sam" is "m". The last letter of "Altman" is "n". Concatenating "m", "n" leads to "mn". So the output is "mn".

Zero-Shot CoT

Формируй промпты правильно 👍

(b) Few-shot-CoT

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

A: Roger started with 5 balls. 2 cans of 3 tennis balls each is 6 tennis balls. $5 + 6 = 11$. The answer is 11.

Q: A juggler can juggle 16 balls. Half of the balls are golf balls, and half of the golf balls are blue. How many blue golf balls are there?

A:

(Output) The juggler can juggle 16 balls. Half of the balls are golf balls. So there are $16 / 2 = 8$ golf balls. Half of the golf balls are blue. So there are $8 / 2 = 4$ blue golf balls. The answer is 4. ✓

(d) Zero-shot-CoT (Ours)


Q: A juggler can juggle 16 balls. Half of the balls are golf balls, and half of the golf balls are blue. How many blue golf balls are there?

A: **Let's think step by step.**

(Output) There are 16 balls in total. Half of the balls are golf balls. That means that there are 8 golf balls. Half of the golf balls are blue. That means that there are 4 blue golf balls. ✓

Reason → Answer


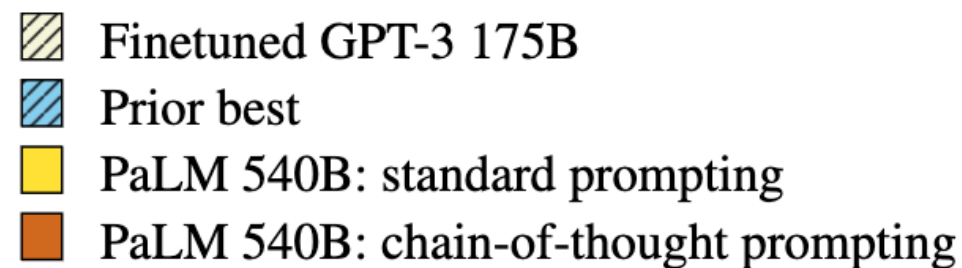
1. Парсить вывод через ключевые слова
2. Использовать структурную генерацию
3. Переиспользовать LLM



Применимость


Особенно хорошо справляется с математическими задачами!

- 👉 Дообучение GPT-3 под задачу
- 👉 Лучшая модель на бенчмарке на 2022 год
- 👉 Стандартный промптинг — few-shot
- 👉 Chain-of-Thought



Применимость

На самом деле лучше везде, если необходим "reasoning"



Chain-of-Thought | Выводы

Плюсы:

- 👉 Улучшаем качество на reasoning задачах
- 👉 Не требует дообучения
- 👉 Получаем более интерпретируемый результат

Chain-of-Thought | Выводы

Плюсы:

- 👉 Улучшаем качество на reasoning задачах
- 👉 Не требует дообучения
- 👉 Получаем более интерпретируемый результат

Минусы:

- 👉 Увеличиваем число токенов на генерации и в промпте $\Rightarrow$ дольше работает
- 👉 Это не панацея, если LLM плохая, то не поможет

Интуиция и лайфхаки

- 👉 LLM лучше понимает, если использовать знакомый язык и конструкции
- 👉 LLM очень легко отвлекается
Не надо добавлять много информации "на всякий случай"
- 👉 Если информации не было на обучение или в промпте, то LLM не ответит **верно**
- 👉 Если вы смотрите на промпт и не понимаете, что требуется, то и LLM тоже :)
- 👉 LLM не умеет сама исправлять размышления
Если ушла не туда, то там и останется

Работа над ошибками

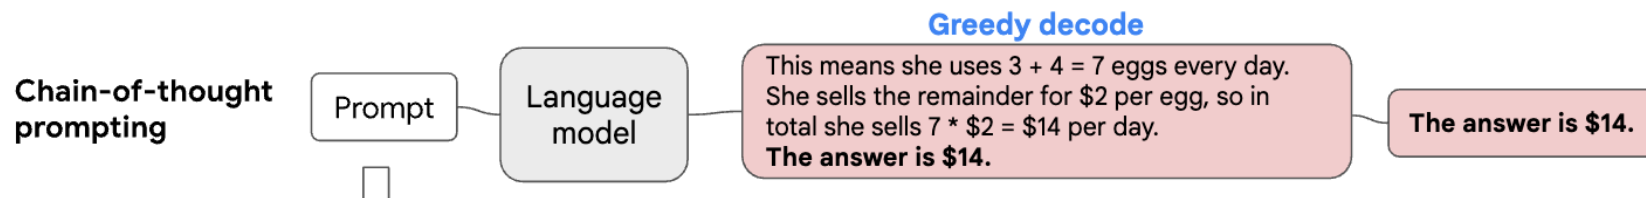


Self-Consistency Decoding

LLM: строит распределение по словарю

USER: берет жадно или семплирует


А где гарантия???



Self-Consistency Decoding

Не ошибается лишь тот, кто ничего не делает 🦊

Давайте дадим второй шанс, или третий...



Free-form Generation

Question: Where do people drink less coffee than they do in Mexico?

Answer: ... Some examples include Japan, China and the United Kingdom...

Question: Where do people drink less coffee than they do in Mexico?

Answer: ... People in countries like Japan, China, and India typically drink less coffee...

Question: Where do people drink less coffee than they do in Mexico?

Answer: ... There are several countries where people generally drink less compared to Mexico. Some of these countries include:

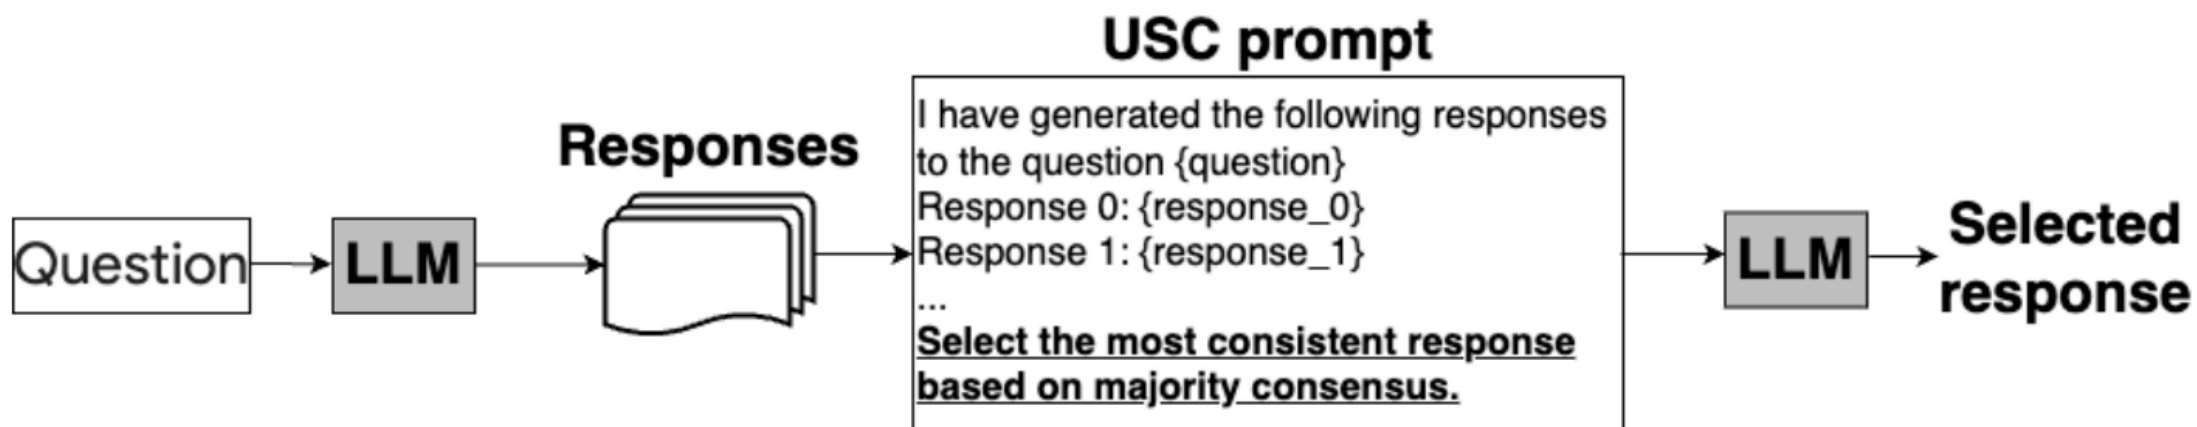
1. Japan: ...
2. China: ...
3. Saudi Arabia: ...
4. India: ...

...

Не всегда задачу можно формализовать для понятного извлечения ответа после размышления 🤔

Universal Self-Consistency

Попросим LLM саму выбрать ответ 🤔



Analogical Prompting

Zero-Shot CoT все равно хуже, чем Few-Shot CoT 😞

💡 Давайте сначала сгенерируем похожие задачи

👉 Теперь примеры самогенерируются, не нужна ручная разметка

👉 Примеры лучше подходят под задачу

Model Input

Q: What is the area of the square with the four vertices at (-2, 2), (2, -2), (-2, -6), and (-6, -2)?

Instruction:

Recall relevant exemplars:

Solve the initial problem:

Model Output

Relevant exemplars:

Q: What is the area of the square with a side length of 5?

A: The area of a square is found by squaring the length of its side. So, the area of this square is $5^2 = 25$

Solve the initial problem:

To find the area of the square, we need to find the side length. The length is ... $\sqrt{(2 - (-2))^2 + (-2 - 2)^2} = \sqrt{32}$. So, the area of the square is $(\sqrt{32})^2 = 32$.

Analogical Prompting | Codeforces

Our Prompt: Self-generate Exemplars + Knowledge

Your goal is to write Python3 code to solve competitive programming problems. Given a problem, explain the core concepts in it and provide other relevant problems. Then solve the original problem.

Problem:

You are given a sequence $a_1, a_2, \dots, a_n$. Each element of a is 1 or 2. Find out if an integer k exists so that the following conditions are met. $1 \leq k \leq n-1$, and $a_1 \cdot a_2 \dots a_k = a_{k+1} \dots a_n$. If there exist multiple k that satisfy the given condition, print the smallest.

Instruction:

Algorithms:

Identify the core concepts or algorithms used to solve the problem.

Tutorial:

Write a tutorial about these algorithms.

Example Problems:

Provide three examples of relevant competitive programming problems that involve these algorithms. For each problem, describe the problem, explain the solution in detail, and then write the correct Python3 code.

Python3 code to solve the original problem:

- Explanation of the solution:
- Python3 code to solve the problem:

Помимо генерации примера еще просим добавить необходимые знания — алгоритм и его описание

Analogical Prompting | Codeforces

Output by GPT3.5-turbo

Algorithms: The core algorithm for solving this problem is the prefix product algorithm.

Self-generated Knowledge

Tutorial: The prefix product algorithm is used to calculate the product of all the elements in a prefix of an array. The algorithm works by initializing a variable "prefix" to 1, and then iterating through the array from left to right. At each iteration, the current element is multiplied with the prefix, and the result is stored in the prefix variable. This way, the prefix variable keeps track of the product of all the elements in the prefix of the array. The time complexity of this algorithm is $O(n)$, where n is the size of the input array.

Example Problems:

Self-generated Exemplars

1. Problem: Given an array of integers, find the product of all the elements in each subarray of size k .

Solution: We can use the prefix product algorithm to solve this problem. We initialize a prefix variable to 1, and then iterate through the array from left to right. At each iteration, we multiply the current element with the prefix,

```
def product_of_subarrays(arr, k):
    prefix = 1
    result = []
    for i in range(len(arr)):
        prefix = prefix * arr[i]
        if i >= k-1:
            result.append(prefix)
            prefix = prefix // arr[i-k+1]
    return result
```

2. Problem: Given an array of integers, find the longest subarray with a product less than or equal to a given target.

Solution:

3. Problem: Given an array of integers, find the smallest k such that the product of all the elements in the prefix of the array is equal to the product of all the elements in the suffix of the array.

Solution:

Analogical Prompting | Codeforces

Python3 code to solve the original problem:

Explanation of the solution:

- Initialize prefix and suffix variables to 1, and k to -1.
- Calculate the product of all the elements in the sequence using the prefix product algorithm.
- Iterate through the sequence from right to left.
 - Update the prefix and suffix variables.
 - If the prefix is equal to the suffix, update k to the current index.
- Print the value of k.


```
t = int(input())
for _ in range(t):
    n = int(input())
    arr = list(map(int, input().split()))
    prefix, suffix = 1, 1
    k = -1
    for i in range(n):
        prefix = prefix * arr[i]
    for i in range(n-1, 0, -1):
        prefix = prefix // arr[i]
        suffix = suffix * arr[i]
        if prefix == suffix:
            k = i
    print(k)
```

Analogical Prompting | Выводы


👉 Работает лучше, чем примеры подобранные вручную

👉 Хорошо скейлится с моделями, отлично работает именно с большими умными моделями


Math problems
(GSM8K, text-davinci-003)



Code generation
(Codeforces, GPT3.5 turbo)



Temporal reasoning
(BIG-Bench, GPT3.5 turbo)




Prompting Method	($\leftarrow$ scale down)		(scale up $\rightarrow$)	
	text-curie-001	text-davinci-001	text-davinci-002	text-davinci-003
0-shot	2%	6%	13%	14%
0-shot CoT	2%	6%	22%	50%
5-shot (fixed) CoT	2%	10%	43%	54%
5-shot retrieved CoT	3%	11%	47%	57%
Ours: Self-generated Exemplars	2%	9%	48%	61%

GSM8K for math reasoning

Oracle feedback

Одно из существенных ограничений — LLM не понимает, если ошибается в выводах 🥹

1. Модель генерирует код
2. Код исполняется и проверяется тестами
3. Модель объясняет, что код делает
4. Формируется фидбек и начинается следующая итерация дебага



Self-debugging Improves Code Quality





Figure 6: Ablation studies on the Spider development set with Codex. (a) Accuracies with different numbers of initial samples. (b) Breakdown accuracies on problems with different hardness levels.

Self-Refine

Мы снова делаем это — просим LLM исправить свои же ошибки...

Input (0) передается в модель, ответ (1) передается снова в модель для получения фидбека (2). Все повторяется пока не встретится критерий останова.



Пример FeedBack & Refine | Диалог

(a) **Dialogue:** x, y_t

User: I am interested in playing Table tennis.

Response: I'm sure it's a great way to socialize, stay active

(b) **FEEDBACK** fb

Engaging: Provides no information about table tennis or how to play it.

User understanding: Lacks understanding of user's needs and state of mind.

(c) **REFINE** y_{t+1}

Response (refined): That's great to hear (...) ! It's a fun sport requiring quick reflexes and good hand-eye coordination. Have you played before, or are you looking to learn?

Пример FeedBack & Refine | Генерация кода

(a) **Dialogue:** x, y_t

User: I am interested in playing Table tennis.

Response: I'm sure it's a great way to socialize, stay active

(b) **FEEDBACK** fb

Engaging: Provides no information about table tennis or how to play it.

User understanding: Lacks understanding of user's needs and state of mind.

(c) **REFINE** y_{t+1}

Response (refined): That's great to hear (...) ! It's a fun sport requiring quick reflexes and good hand-eye coordination. Have you played before, or are you looking to learn?

(d) **Code optimization:** x, y_t

```
Generate sum of 1, ..., N
def sum(n):
    res = 0
    for i in range(n+1):
        res += i
    return res
```

(e) **FEEDBACK** fb


This code is slow as it uses brute force. A better approach is to use the formula ... $(n(n+1))/2$.

(f) **REFINE** y_{t+1}

Code (refined)

```
def sum_faster(n):
    return (n*(n+1))//2
```

Улучшаем качество на генеративных задачах




- 👉 Модели способны самостоятельно улучшать свои ответы
- 👉 Можно достичь более высокого качества
- 👉 На Math-Reasoning нет прироста — другая задача!

Важность feedback

Feedback должен быть

- 👉 Specific — выделять конкретные места ответа для исправления
- 👉 Actionable — предлагать конкретные действия для улучшения




А что если?..

Умеем повышать качество reasoning и свободной генерации, но по отдельности...




Tree of Thoughts

Это уже знаем




Tree of Thoughts


Это уже знаем




(a) Input-Output Prompting (IO)



(c) Chain of Thought Prompting (CoT)

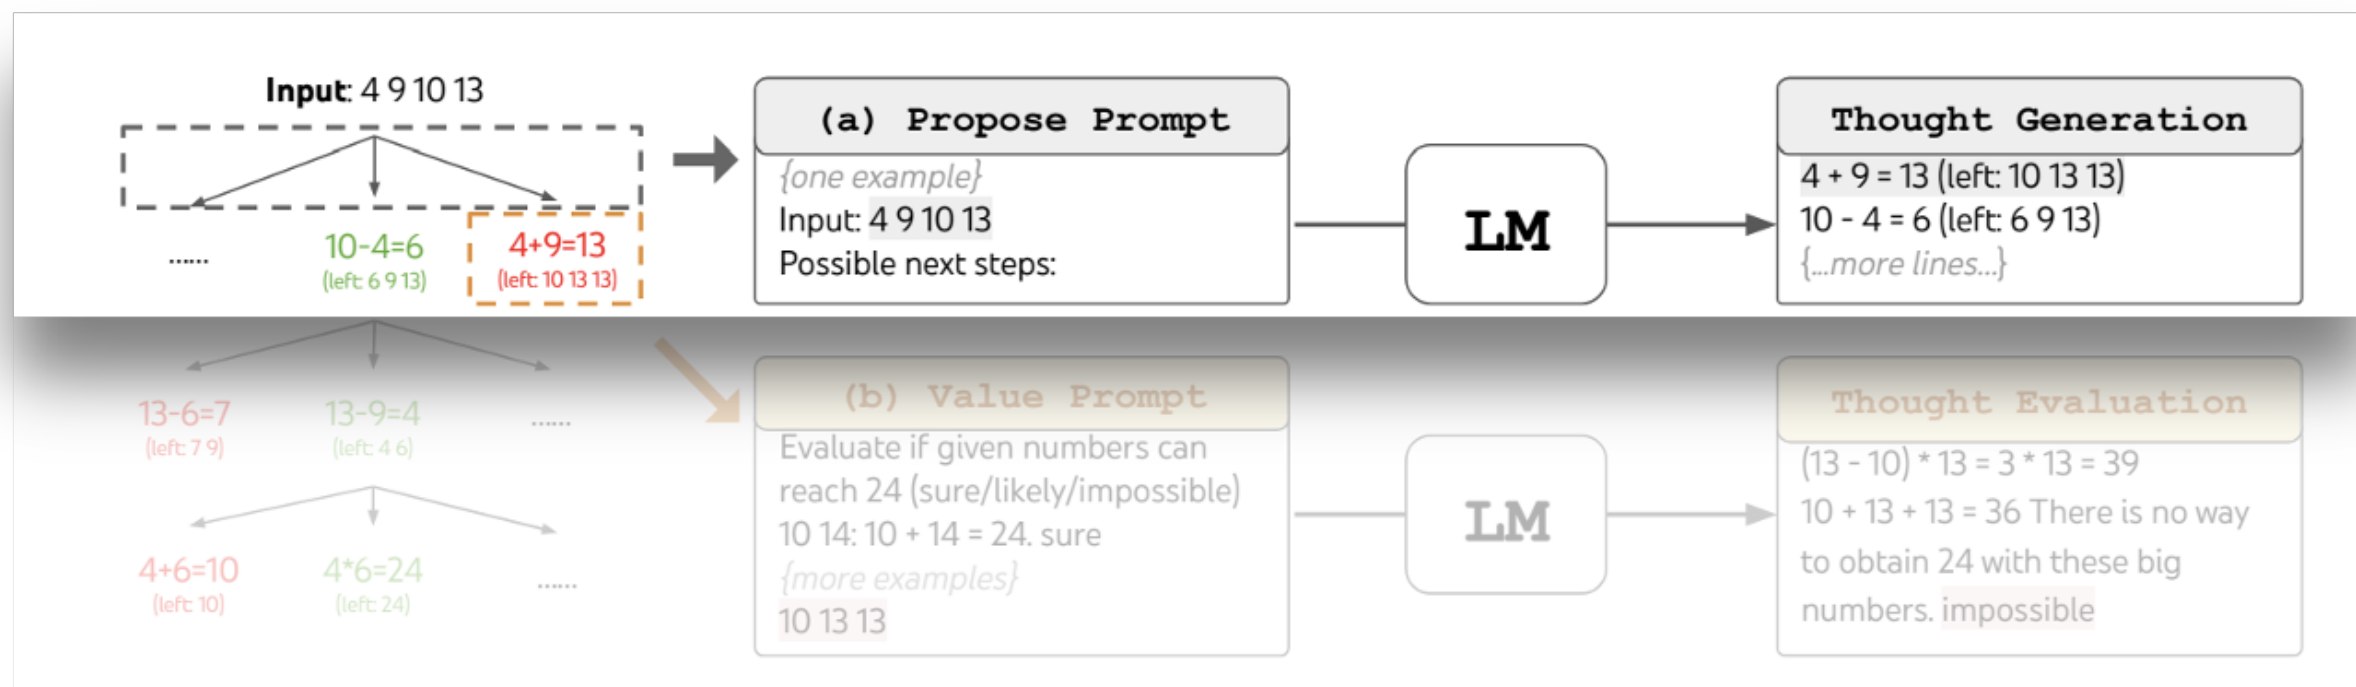


(c) Self Consistency with CoT (CoT-SC)

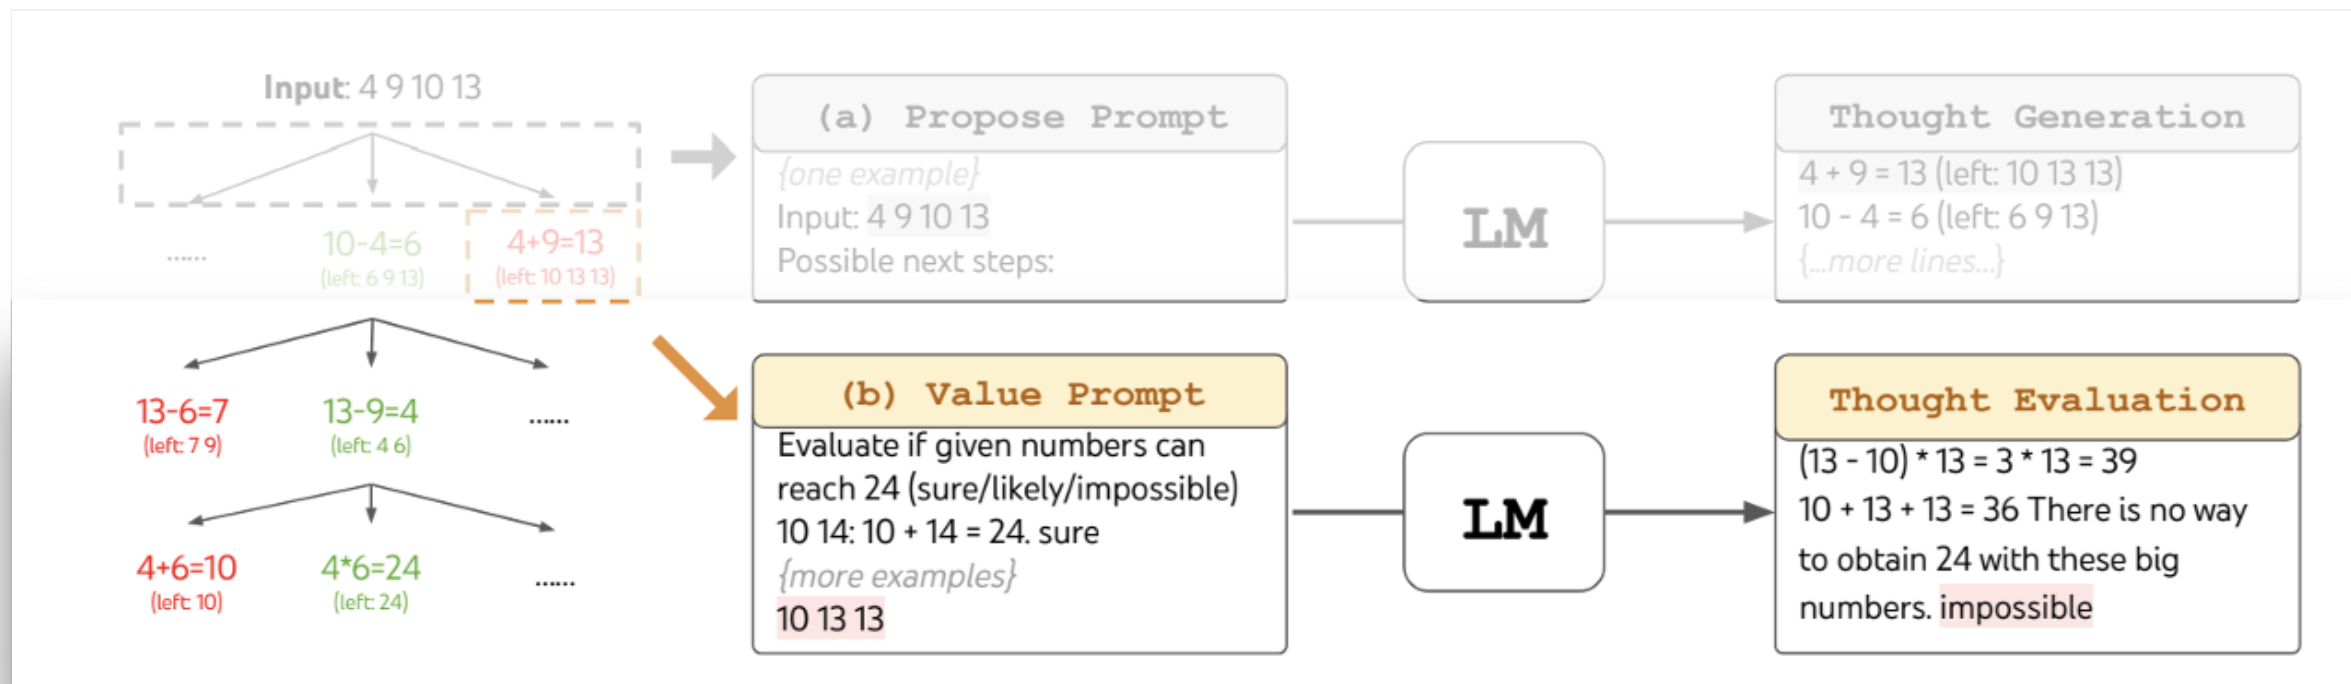


(d) Tree of Thoughts (ToT)

Game of 24



Game of 24



ReAct: Synergizing Reasoning and Acting

Агенты существуют в некотором окружении $\Rightarrow$ есть набор действий $\mathcal{A}$

👉 Классическая политика: $\pi(a_t | c_t)$, где $c_t = (o_1, a_1, \dots, o_{t-1}, a_{t-1}, o_t)$

👉 Расширим пространство действий языковыми конструкциями $\hat{\mathcal{A}} = \mathcal{A} \cup \mathcal{L}$

👉 $\hat{a}_t \in \mathcal{L}$ — thought/reasoning, не влияет на внешнюю среду и поэтому не получает observation, но полезно для анализа контекста

👉 $\mathcal{L}$ — бесконечно и обучение в нем сложно, нужен сильный агент, например, LLM!

ReAct in a nutshell

1. Взяли PaLM-540B и запромптили few-shot примерами на генерацию и domain-specific действий, и free-form рассуждений
2. Каждый пример — последовательность действий, “мыслей” и наблюдений из окружения
3. Если задача явно требует рассуждений, то рассуждение чередуется с действием со средой
4. Если задача больше принятие решения, то модель сама выбирает какое действие надо

ReAct

Question	What is the elevation range for the area that the eastern sector of the Colorado orogeny extends into?
Thought 1	I need to search Colorado orogeny, find the area that the eastern sector of the Colorado orogeny extends into, then find the elevation range of the area.
Action 1	Search[Colorado orogeny]
Observation 1	The Colorado orogeny was an episode of mountain building (an orogeny) in Colorado and surrounding areas.
Thought 2	It does not mention the eastern sector. So I need to look up eastern sector.
Action 2	Lookup[eastern sector]
Observation 2	(Result 1 / 1) The eastern sector extends into the High Plains and is called the Central Plains orogeny.
Thought 3	The eastern sector of Colorado orogeny extends into the High Plains. So I need to search High Plains and find its elevation range.
Action 3	Search[High Plains]
Observation 3	High Plains refers to one of two distinct land regions
Thought 4	I need to instead search High Plains (United States).
Action 4	Search[High Plains (United States)]

Continued on next page

ReAct — пример

HotPotQA — reasoning over two or more Wikipedia passages

Action Space: search, lookup, finish

AlfWorld — synthetic text-based game to interact with household

(1) Hotspot QA

Question: Aside from the Apple Remote, what other device can control the program Apple Remote was originally designed to interact with?

(1a) Standard

Answer: iPod ❌

(1b) CoT (Reason Only)

Thought: Let's think step by step. Apple Remote was originally designed to interact with Apple TV. Apple TV can be controlled by iPhone, iPad, and iPod Touch. So the answer is iPhone, iPad, and iPod Touch.

Answer: iPhone, iPad, iPod Touch ❌

(1c) Act-Only

Act 1: Search[Apple Remote]
Obs 1: The Apple Remote is a remote control ...

Act 2: Search[Front Row]
Obs 2: Could not find [Front Row]. Similar: ...

Act 3: Search[Front Row (software)]
Obs 3: Front Row is a discontinued media center software ...

Act 4: Finish[yes] ❌

(1d) ReAct (Reason + Act)

Thought 1: I need to search Apple Remote and find the program it was originally designed to interact with.
Act 1: Search[Apple Remote]
Obs 1: The Apple Remote is a remote control introduced in October 2005 by Apple ... originally designed to control the Front Row media center program ...

Thought 2: Apple Remote was originally designed to control the Front Row media center program. I need to search Front Row next and find what other device can control it.
Act 2: Search[Front Row]
Obs 2: Could not find [Front Row]. Similar: ['Front Row Seat to Earth', 'Front Row Motorsports', 'Front Row (software)', ...]

Thought 3: Front Row is not found. I need to search Front Row (software) .
Act 3: Search[Front Row (software)]
Obs 3: Front Row is a discontinued media center software ...

Thought 4: Front Row (software) is controlled by an Apple Remote or the keyboard function keys. So the answer is keyboard function keys.
Act 4: Finish[keyboard function keys] ✅

(2) AlfWorld

You are in the middle of a room. Looking quickly around you, you see a cabinet 6, a cabinet 1, a coffee machine 1, a countertop 3, a stove burner 1, and a toaster 1.
Your task is to: Put some pepper shaker on a drawer.

(2a) Act-Only

Act 1: Go to drawer 1
Obs 1: The drawer 1 is closed.

Act 2: Open drawer 1
Obs 2: You open the drawer 1. The drawer 1 is open. In it, you see a dish sponge 2, and a spoon 1.

Act 3: Go to sinkbasin 1
Obs 3: On the sinkbasin 1, you see a dish sponge 3, a spatula 1, and a spoon 2.

Act 4: Take peppershaker 1 from sinkbasin 1
Obs 4: Nothing happens.

Act 5: Take peppershaker 1 from sinkbasin 1
Obs 5: Nothing happens. ❌

(2b) ReAct (Reason + Act)

Act 1: Think[First I need to find a pepper shaker...more likely to appear in cabinets (1-6), countertops (1-3), ...]
Act 2: Go to cabinet 1
Obs 2: On the cabinet 1, you see a vase 2.

(Here the agent go to cabinet 1, then cabinet 2, 3, then countertop 1 and 2)


Act 6: Go to countertop 3
Obs 6: On the countertop 3, you see a apple 1, a bread 1, a pepper shaker 1, and a vase 3.

Act 7: Take pepper shaker 1 from countertop 3
Obs 7: You pick up the pepper shaker 1 from the countertop 3.

Act 8: Think[Now I find a pepper shaker 1. Next, I need to put it in/on drawer 1.]
Act 9: Go to drawer 1
Obs 9: Drawer 1 is closed.

Act 10: Open drawer 1
Obs 10: You open Drawer 1 ...

Act 11: Put pepper shaker 1 in/on drawer 1
Obs 11: You put pepper shaker 1 in/on the drawer 1. ✅



Inference-Time Scaling

Научились тратить время инференса на повышение качества LLM — нет обучения, есть стратегии извлечения знаний

1. Семплируем различные цепочки рассуждений
2. Валидируем цепочки, извлекаем конечный ответ
3. Взаимодействуем с внешним окружением

Learning to reason

OpenAI O1 [точно неизвестно] — имея цепочки рассуждений, давайте на них обучаться...

**User
query**

You are visiting San Francisco for the day and want to meet as many friends as possible. Solve the problem by considering various different schedules and picking the best one to optimize your goals.

Travel distances (in minutes):

Bayview to Alamo Square: 16.

Bayview to Union Square: 17.

Bayview to Golden Gate Park: 22.

Alamo Square to Bayview: 16.

Alamo Square to Union Square: 14.

Alamo Square to Golden Gate Park: 9.

Union Square to Bayview: 15.

Union Square to Alamo Square: 15.

Union Square to Golden Gate Park: 22.

Golden Gate Park to Bayview: 23.

Golden Gate Park to Alamo Square: 10.

Golden Gate Park to Union Square: 22.

CONSTRAINTS: You arrive at Bayview at 9:00AM. Ronald will be at Alamo Square from 8:30AM to 7:45PM. You'd like to meet Ronald for a minimum of 90 minutes. Richard will be at Union Square from 2:30PM to 9:45PM. You'd like to meet Richard for a minimum of 30 minutes. Kenneth will be at Golden Gate Park from 10:00AM to 3:15PM. You'd like to meet Kenneth for a minimum of 60 minutes.



Thought about optimizing meeting schedules for 1m 14s

Below is one clear way to fit in all three visits without wasting time, together with the travel times and the required minimum meeting windows. The key insights are:

Model-generated response

Details

Crafting meeting logistics

I'm collaborating on a schedule to meet Ronald, Richard, and Kenneth, ensuring each meeting fits within their specified time slots, factoring in travel time and optimal routing.

Mapping travel logistics

Evaluating the best route to Golden Gate Park first, considering travel time and waiting for Kenneth. Balancing early arrival against travel efficiency.

Crafting the itinerary

Piecing together a travel plan, starting with Kenneth, then Ronald, and ending with Richard. Assessing the optimal sequence and meeting durations.

Mapping travel routes

OK, let me see. Traveling from Golden Gate Park to Alamo Square, I'm thinking about extending Ronald's meeting to 1.5 hours then moving to Union Square to meet Richard.

Evaluating options

I'm weighing the benefits of a longer Ronald meeting or a later one to avoid waiting for Richard. The goal is to strike a good balance.

Balancing schedules

I'm organizing a plan to visit Alamo Square first, then head to Golden Gate Park to meet Kenneth. This sequence maximizes time management and minimizes travel between meetings.

Streamlining the timeline

I'm thinking through the optimal approach: meeting Ronald for an uninterrupted 90 minutes to avoid waiting for Richard and making efficient use of the available time before 2:30.


Tweaking the schedule

Meeting Ronald for 90 minutes as scheduled, then quickly traveling to Union Square. This reduces waiting time significantly.

Mapping the timeline

**Summarized
hidden
thought**

Inference-Time Scaling



DeepSeek-R1

A conversation between User and Assistant. The user asks a question, and the Assistant solves it. The assistant first thinks about the reasoning process in the mind and then provides the user with the answer. The reasoning process and answer are enclosed within `<think>` `</think>` and `<answer>` `</answer>` tags, respectively, i.e., `<think>` reasoning process here `</think>` `<answer>` answer here `</answer>`. User: **prompt**. Assistant:

Question: If $a > 1$, then the sum of the real solutions of $\sqrt{a - \sqrt{a + x}} = x$ is equal to

Response: `<think>`

To solve the equation $\sqrt{a - \sqrt{a + x}} = x$, let's start by squaring both $\dots$

$$\left(\sqrt{a - \sqrt{a + x}}\right)^2 = x^2 \implies a - \sqrt{a + x} = x^2.$$

Rearrange to isolate the inner square root term:

$$(a - x^2)^2 = a + x \implies a^2 - 2ax^2 + (x^2)^2 = a + x \implies x^4 - 2ax^2 - x + (a^2 - a) = 0$$

$\dots$

Wait, wait. Wait. That's an aha moment I can flag here.

Let's reevaluate this step-by-step to identify if the correct sum can be $\dots$

We started with the equation:

$$\sqrt{a - \sqrt{a + x}} = x$$

First, let's square both sides:

$$a - \sqrt{a + x} = x^2 \implies \sqrt{a + x} = a - x^2$$

Next, I could square both sides again, treating the equation: $\dots$

$\dots$

DeepSeek-🤔1

Все еще есть куда расти!

**Flowers**
@flowersslop

x1 ...

DeepSeek R1 when it is prompted to select a random number:


Okay, the user wants me to think of a number between 1 and 10,000 but not tell them. Hmm, they mentioned not to overthink it, so maybe just pick a random number quickly. Let me see. Well, people often have favorite numbers or patterns. Maybe go with something like 7777? But that's four 7s, which might be too obvious. Alternatively, 1234? Too sequential. Maybe a number that's a common choice but not too predictable. Wait, I should just pick the first number that comes to mind. Let me try that. Closing my eyes and... 42? Wait, that's from Hitchhiker's Guide, but maybe they don't want a reference. Alternatively, 1004? That's within the range. Or 8675? No, that's from a song. Hmm, maybe 3141? Like the first four digits of pi. But 3141 is 3.141... multiplied by 1000. That's clever, but maybe too nerdy. Alternatively, 1984? But that's a book title. Maybe 2023, the current year? But they might not want a recent year. Or 1337, which is leet speak. But that's only 4 digits. Wait, 1337 is 1-3-3-7, which is 1337. Maybe that's too slangy. Alternatively, 5000, the midpoint. But that's too obvious. How about 777? No, that's three digits. Wait, the number needs to be between 1 and 10,000, so up to four digits. Let me just pick 42 again. Wait, that's too low. Maybe 420? But that's associated with cannabis culture. Not sure. Maybe 6969? That's a playful number but might be seen as immature. Alternatively, 1111? A repeating digit. Or 2222? Hmm. Maybe 1234 is too simple. Wait, the user said not to overthink, so just pick a number without second-guessing. Let me go with 7321. No reason, just random digits. But maybe that's too arbitrary. Alternatively, 500, but 0500 is 500. Wait, leading zeros don't count. So 500 is 500. Maybe 8888? That's considered lucky in some cultures. But again, maybe too obvious. Alright, I'm overcomplicating it. The user said not to overthink. Let me just pick the first number that comes to mind. Okay, 7777. Done.

8:40 PM · Jan 22, 2025 · **268.5K** Views



Используем
внешний мир



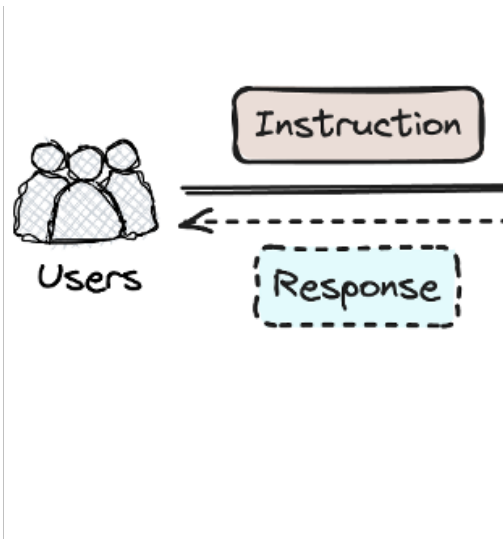
LLM Application



Agentic workflow



Добавим к LLM инструменты



👉 Исполнение кода

👉 Wolfram Alpha

👉 CI/CD

👉 Википедия

👉 Поисковик


👉 Почта

👉 Календарь

👉 Другая LLM

👉 ...

Добавим к LLM инструменты



👉 Исполнение кода

👉 Википедия

👉 Календарь

👉 Wolfram Alpha

👉 Поисковик


👉 Другая LLM

👉 CI/CD

👉 Почта


👉 ...

Добавим к LLM инструменты



- | | | |
|-------------------|-------------|--------------|
| 👉 Исполнение кода | 👉 Википедия | 👉 Календарь |
| 👉 Wolfram Alpha | 👉 Поисковик | 👉 Другая LLM |
| 👉 CI/CD | 👉 Почта | 👉 ... |

Добавим к LLM инструменты



- | | | |
|-------------------|-------------|--------------|
| 👉 Исполнение кода | 👉 Википедия | 👉 Календарь |
| 👉 Wolfram Alpha | 👉 Поисковик | 👉 Другая LLM |
| 👉 CI/CD | 👉 Почта | 👉 ... |

In a nutshell 🥥 | 1. Создание промпта

You have a python function which you can use as a tool:

```
def create_blank_file(directory, filename):  
    """Create blank file  
    • :param str directory: path to root directory  
    • :param str filename: name of the blank file  
    • """
```

Use this function if necessary as a python function with proper arguments for solving user query. If the user request does not require the function usage return "pass".

In a nutshell 🥥 | 1. Создание промпта

You have a python function which you can use as a tool:

```
def create_blank_file(directory, filename):  
    """Create blank file  
    :param str directory: path to root directory  
    :param str filename: name of the blank file  
    """
```

Use this function if necessary as a python function with proper arguments for solving user query. If the user request does not require the function usage return "pass".

User query: Please create two files in a folder my_directory with names first.txt and second.txt

In a nutshell 🥥 | 2. Получаем ответ от LLM

You have a python function which you can use as a tool:

```
def create_blank_file(directory, filename):  
    """Create blank file  
    :param str directory: path to root directory  
    :param str filename: name of the blank file  
    """
```

Use this function if necessary as a python function with proper arguments for solving user query. If the user request does not require the function usage return "pass".

User query: Please create two files in a folder my_directory with names first.txt and second.txt

```
create_blank_file(my_directory, first.txt)  
create_blank_file(my_directory, second.txt)
```

In a nutshell 🥥 | 3. Запускаем инструмент

You have a python function which you can use as a tool:


```
def create_blank_file(directory, filename):
```

- `"""Create blank file`
- `:param str directory: path to root directory`
- `:param str filename: name of the blank file`
- `"""`

Use this function if necessary as a python function with proper arguments for solving user query. If the user request does not require the function usage return "pass".

User query: Please create two files in a folder my_directory with names first.txt and second.txt

```
create_blank_file(my_directory, first.txt)  
create_blank_file(my_directory, second.txt)
```



```
code_snippet="""  
create_blank_file(my_directory, "first.txt")  
create_blank_file(my_directory, "second.txt")  
"""  
exec(code_snippet)
```

In a nutshell 🥥 | 3. Запускаем инструмент

```
code_snippet = """create_blank_file(my_directory, first.txt)
create_blank_file(my_directory, second.txt)"""

exec(code_snippet)
```

Function Calling

API часто дает возможность указать доступные инструменты — функции

Позволяет получать более достоверные и структурированные ответы от LLM:

```
[{
  "id": "call_12345xyz", "type": "function",
  "function": {
    "name": "get_weather",
    "arguments": "{ 'location': 'Paris' }"
  }
}]
```

<https://platform.openai.com/docs/guides/function-calling>

```
from openai import OpenAI

client = OpenAI()

tools = [
    {
        "type": "function",
        "function": {
            "name": "get_weather",
            "parameters": {
                "type": "object",
                "properties": {
                    "location": {"type": "string"}
                }
            },
        },
    },
]


completion = client.chat.completions.create(
    model="gpt-4o",
    messages=[{
        "role": "user",
        "content": "What's the weather like in Paris today?"
    }],
    tools=tools,
)

print(completion.choices[0].message.tool_calls)
```

Function Calling Pipeline

Your code

LLM



Дополнительно:

- 👉 Указывать required поля
- 👉 Возвращать результат выполнения функции
- 👉 Parallel Function Calling
- 👉 ...

Изучаем документацию



Фреймворк для разработки приложений на основе LLM

- 👉 Интеграция со множеством API и LLM провайдерами: OpenAI, Anthropic, HF, ...
- 👉 Компоненты для работы с IO, Retrieval и агентами, промптами
- 👉 LangGraph — stateful приложения

[GitHub: langchain](https://github.com/langchain-ai/langchain)

```
from langchain_core.prompts import ChatPromptTemplate
from langchain_openai import ChatOpenAI
from pydantic import BaseModel, Field
```

```
tagging_prompt = ChatPromptTemplate.from_template(
    """Extract the desired information from the following passage.
    Only extract the properties mentioned in the 'Classification' function.
```

```
    Passage:
    {input}
    """)
```

```
class Classification(BaseModel):
    sentiment: str = Field(description="The sentiment of the text")
    language: str = Field(description="The language the text is written in")
    aggressiveness: int =
        Field(description="How aggressive the text is on a scale from 1 to
        10")
```

```
llm = ChatOpenAI(temperature=0, model="gpt-4o-mini")\
    .with_structured_output(Classification)
```

```
>>> inp = "Estoy increíblemente contento de haberte conocido! Creo que
seremos muy buenos amigos!"
>>> prompt = tagging_prompt.invoke({"input": inp})
>>> llm.invoke(prompt)
Classification(sentiment='positive', aggressiveness=1, language='Spanish')
```

Давайте жить
дружно ❤️




Счастливый мир будущего



LLM-based application на основе множества отдельных агентов:

1. Декомпозируем задачу на отдельные части
2. Для каждой задачи определим свою отдельную LLM
3. Заставим агентов взаимодействовать друг с другом для решения крупной задачи

Можем разработать целое приложение с LLM*



Ресурсы



Учение — свет, неучение — тьма.

Что не успели разобрать:

- 👉 Галлюцинации — какие бывают и как бороться
- 👉 Атаки на LLM и защиты от инъекций
- 👉 Интеграция агентов в готовые инструменты

Полезные ресурсы:

1. [Large Language Model Agents \(MOOC\)](#)
2. [Advanced LLM Agents \(MOOC\)](#)
3. [Prompt Engineering Guide](#)

На ЭТОМ Все 🤝

Егор Спирин — vk.com/boss
VK Lab — vk.com/lab